

A Brief Survey of Multi-Processor Scheduling For Hard Real-Time Systems

Xin Lin^a, Xiaorong Zhu^a, Lijia Liu^a

^a*Department of Computer Science, The University of Texas at Austin*

Abstract

In class, both of scheduling algorithms [1] and priority inheritance protocols [2] in the context of a single processor were examined in details. Nevertheless, the emergence and popularity of distributed computing system gave rise to the need to solve multi-processor scheduling and priority inheritance problems. As the supplementary study, this paper surveys existing scheduling algorithms in the context of multiple processors. The very first section outlines the background of multi-processor scheduling problems, as well as system models, terminology, and the metrics of scheduling algorithms. After that, partitioned scheduling and global scheduling, as the primary objects of our research, will be fully explored. Moreover, we will also give brief sketch to the hybrid approaches of partitioned scheduling and global scheduling.

Keywords: System, Scheduling Algorithm, Task Management

1. Introduction

Real-time systems has become prevalent in diverse application areas, such as traffic controls, embedded automotive electronics, networked multimedia etc. Typically, a real-time system comprises a real-time computer system and a controlled object, which are connected through sensors, actuators and some other input-output interfaces. The controlling real-time computer reacts/responds to its controlled object based on the currently known information about the object. For example, a real-time computer controls a device according to the data read by sensors at periodic intervals and responds by sending signals to actuators within a time bound. (Figure?)

The purpose of real-time systems is to meet various timing constraints imposed on it by its external world during the operation of the system. The

13 instant at which a response from the controlling system must be delivered to
14 its controlled system is called a deadline. A deadline can be classified into
15 three categories[3]:

- 16 • Soft Deadline: the response still has effectiveness even after the deadline
17 has passed. However the later the execution of a task after the deadline,
18 the more penalty it pays.
- 19 • Firm Deadline: the response has no utilize if the deadline has passed.
20 Moreover the earlier the execution of a task before a firm deadline, the
21 more rewards it gains.
- 22 • Hard Deadline: a task must complete its computation by the deadline,
23 otherwise a catastrophe could happen.

24 Commands and Control systems, Air traffic control systems are examples
25 for hard real-time systems. On-line transaction systems, airline reservation
26 systems are soft real-time systems[4]. Such a system is called embedded
27 real-time system if it is a component of a larger system.

28 Companies are motivated to build embedded real-time systems by their
29 advanced effectiveness and flexibility. In order to beat the competition in
30 industries, they need to satisfy customers needs and lower associated costs
31 to the maximum extent possible. Such aspiration has resulted in a rapid
32 progress in the area of embedded real-time technology. For instance, sili-
33 con vendors used to increase processor performance by concentrating on the
34 miniaturization needed. However this approach has led to problems with
35 both high power consumption and excessive heat dissipation[4]. Thus using
36 multiprocessor platforms for high-end real-time applications is an increasing
37 trend toward instead[4].

38 Multiprocessor systems can be classified into three categories from the
39 perspectives of scheduling[4]:

- 40 • Heterogeneous: processors in the system are different. Not all tasks
41 may be able to execute on all processors and the execution rate of a
42 task depends on both the processor and the task.
- 43 • Homogeneous: processors in the system are identical. The execution
44 rate of all tasks is the same on all processors. Thus it depends only on
45 the task.

- Uniform: the execution rate of a task depends only on the speed of processors. For example a processor with speed 2 will execute all tasks at exactly twice the rate of a processor with speed 1.

In this paper we are concerned with scheduling algorithms of **homogeneous** multiprocessor for **hard** real-time system. The paper is organized as follows: section 2 is problem statement, section 3 is key concepts, section 4 is performance evaluation. we will describe the partitioned scheduling and global scheduling in detail in section 5 and section 6, respectively. A conclusion is given at the end of this paper (need to revise).

2. Preliminaries

This section reviews the fundamental terminologies and notations used in multiprocessor scheduling algorithms.

2.1. Multiprocessor Scheduling

In computer science, multiprocessor scheduling is an NP-hard optimization problem[wiki]. The problem can be described as "given a set J of jobs where job j_i has length l_i and a number of processors m , what is the minimum possible time required to schedule all jobs in J on m processors such that none overlap?"[5]. Clearly multiprocessor schedulers should concern with the allocation of the resources and the satisfaction of timing constraints[4]:

- Resource Allocation: on which processor a task executes.
- Propriety Determination: when and in what order, with respect to jobs of other tasks, each job executes so that all jobs can be finished before their deadlines.

Then real-time scheduling algorithms for multiprocessor systems can be classified from perspectives of allocation changes and priority changes (referred to as migration-based and priority-based classifications [6]).

2.2. Classification of Real-time Scheduling Algorithms

Real-Time scheduling can be categorized into hard scheduling and soft scheduling based on the types of deadlines. The hard real-time scheduling can be further classified into two types: static and dynamic. A static scheduling generates its scheduling decisions off-line according to the prior

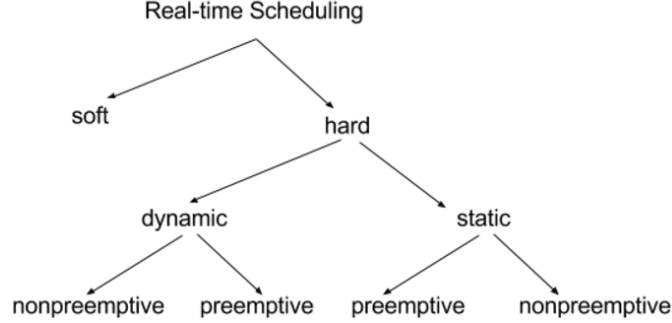


Figure 1: Taxonomy of Real-time Scheduling Algorithms

information of the jobs in a task. The decisions are made at compile time, thus its run-time overhead is small. By contrast, the scheduling decisions are made at run time in dynamic scheduling system, which leads to significant overheads. Furthermore, both dynamic scheduling and static scheduling can have preemptive or nonpreemptive scheduling:

- Preemptive: the currently executing task can be preempted by a higher priority task at anytime
- Nonpreemptive: the currently executing task can not be preempted until completion

In this paper, we are mainly concerned with dynamic preemptive scheduling algorithms. Figure 1 illustrates a rough taxonomy of real-time scheduling algorithms.

More specifically, the dynamic scheduling algorithm can be classified by its priority policy[4]:

- Fixed task-level priority: the jobs of a task has a single fixed priority, such as Rate Monotonic RM algorithm
- Fixed job-level priority: the jobs of a task may have different priorities, but each job has a single static priority, such earliest deadline first (EDF) scheduling.
- Dynamic priority: A single job may have different priorities at different times, such as least laxity first (LLF) scheduling.

98 From the respective of allocation policy, the real-time scheduling algo-
99 rithms for multiprocessor systems can be also classified as follows[4]:

- 100 • Non-migration: the jobs of a task must execute on one single processor
101 and migration is not permitted.
- 102 • Task-level migration: the jobs of a task may be allocated to different
103 processors, however, each job can only execute on one single processor.
- 104 • Job-level migration: a single job can be executed on different proces-
105 sors, however, no parallel execution of a job is permitted.

106 Scheduling algorithms without migrations are referred partitioned algo-
107 rithms while those that permits migrations are referred as global algorithms.
108 These two algorithms have been discussed in detail in section 3 and section
109 4, respectively.

110 3. Partitioned Scheduling

111 In this section, we will review some partitioned approaches to multipro-
112 cessor real-time scheduling, and then compare their performance.

113 3.1. Characteristic of Partitioned Scheduling

114 Partitioned scheduling provides capability for performing parallel pro-
115 cessing and also for automation of batch execution of multiple processes. It
116 can improve performance by breaking down a process that works on a large
117 data set, to multiple parallel processes that work on smaller data sets, each
118 contains part of the original data set. The general solution involves two algo-
119 rithms: one to assign tasks to processors, known as the allocation, the other
120 to schedule tasks that are assigned to each individual processor. The parti-
121 tioning strategy also requires that all occurrences of a task to be executed
122 on the same processor.

123 One main advantage of partitioned scheduling is that, after the allocation
124 of tasks to processors, we can apply the optimal real-time scheduling tech-
125 niques and analyses for uni-processor systems to each individual processor.

126 Comparing to the global scheduling, partitioned scheduling also has some
127 advantages.

- 128 • A task that overruns its worst-case execution time will not affect tasks
129 on other processors.

- 130 • There is no migration cost as each task only runs on a single processor.
- 131 • For each processor, the run-queue is much smaller comparing to the
- 132 single global queue in the global scheduling, thus the overheads of man-
- 133 aging the run-queue can be neglected.

134 However, the main disadvantage of the partitioned approach is that the
 135 allocation problem is analogous to a NP-Hard problem.

136 3.2. Rate Monotonic Next Fit Scheduling (RMNF)

137 The rate-monotonic-next-fit scheduling [7] is a partitioned scheduling al-
 138 gorithm for multiprocessors systems. According to this algorithm, tasks are
 139 first sorted in the descending order of their request rates. The scheduling
 140 scheme is that:

- 141 1. Starts from the first task i , and the first processor j .
- 142 2. Assign the task i to the processor j if j meets the requirement that
- 143 together with all tasks that have been assigned to the processor j , tasks
- 144 can be feasibly scheduled on processor j according to the rate-monotonic
- 145 scheduling algorithm for a single processor. Otherwise, assign i to $j +$
- 146 1 and increase j by 1.
- 147 3. Get the next task and repeat the previous two steps, unless there is no
- 148 task left.

149 Let N be the number of processors required to be feasibly schedule a set of
 150 task by the RMNF algorithm, and N_0 be the minimum number of processors
 151 required to feasibly schedule the same set of tasks. Then as N_0 approaches
 152 infinity, we can get that [7]

$$2.4 \leq \frac{N}{N_0} \leq 2.67 \quad (1)$$

153 3.3. Rate Monotonic First-Fit Scheduling (RMFF)

154 The rate-monotonic-first-fit scheduling algorithm [7] is similar to the
 155 RMNF. According to the RMFF, the tasks are first sorted in the descending
 156 order of their request rates. We will use a counter N to be the number of
 157 processors required for scheduling the given tasks. N is first initiated to 1.
 158 The scheduling scheme is that:

- 159 1. Starts from the first task i .

- 160 2. Starts from the lowest-indexed processor. Assign the task i to the first
 161 processor that meets the requirement that together with all tasks that
 162 have been assigned to the processor, tasks can be feasibly scheduled
 163 on the processor according to the rate-monotonic scheduling algorithm
 164 for a single processor. If no processor with index less than N meets the
 165 requirement, increase N by 1.
- 166 3. Get the next task and repeat the previous step, unless there is no task
 167 left.

168 Similarly, the bound of the rate of the number of processors N required to
 169 be feasibly schedule the task set by this algorithm, and the minimum number
 170 of processors N_0 required to be feasibly schedule the same set of tasks can
 171 be obtained when N_0 approaches infinity:

$$2 \leq \lim_{N_0 \rightarrow \infty} \frac{N}{N_0} \leq \frac{4 \times 2^{1/3}}{1 + 2^{1/3}} \approx 2.33 \quad (2)$$

172 3.4. Rate Monotonic Best-Fit Scheduling (RMBF)

173 The rate-monotonic-best-fit scheduling [8] algorithm is based on the bin-
 174 packing heuristic. Best-Fit scheme chooses to assign tasks on a processor
 175 that can maximize the utilization of that processor.

176 Similar to RMNF and RMFF, we first sort all given tasks according to
 177 non-decreasing periods. Keep a counter N to be the number of processors
 178 required for scheduling the given set of tasks. N is set to 1 at the beginning.
 179 The scheduling algorithm works like this:

- 180 1. Starts from the first task i .
- 181 2. Starts from the lowest-indexed processor j . Let k_j and U_j denote the
 182 number of tasks already assigned to the processor j and the total uti-
 183 lization of the k_j tasks. Let u_i denote the utilization of the task i . Find
 184 the smallest-indexed processor j such that task i together with all k_j
 185 tasks can be feasibly scheduled by the rate-monotonic scheduling algo-
 186 rithm, and $2(1 + \frac{U_j}{k_j})^{-k_j} - 1$ be as small as possible, then assign task i
 187 to j , and set $k_j = k_j + 1$, $U_j = u_i + U_j$. If $j < m$ then set $m = j$.
- 188 3. Get the next task and repeat the previous step, unless there is no task
 189 left.

190 Similarly, the bound of rate of N and N_0 (with the same definitions in
 191 RMFF) can be obtained when N_0 approaches infinity:

$$2.3 \leq \lim_{N_0 \rightarrow \infty} \frac{N}{N_0} \leq 2 + \frac{3 - 2^{3/2}}{2(2^{1/3} - 1)} \approx 2.33 \quad (3)$$

Table 1: Approximation Ratios

Algorithm	Approximation Ratio $\frac{N}{N_0}$
RMNF	2.67
RMFF	2.33
RMBF	2.33

192 3.5. Comparison and Summary

193 In this section, we had introduced the three partitioned scheduling algo-
194 rithms RMNF, RMFF and RMBF. All these three algorithms require that
195 tasks are sorted according to their non-decreasing periods/request rates.
196 They are only applicable to be applied to the situations that the periods
197 of incoming tasks are fixed and static. The performance of these three algo-
198 rithms are evaluated in the approximation ratio, which is defined to be the
199 rate of N and N_0 , the number of processors required to be feasibly sched-
200 ule a set of task by the algorithm and the minimum number of processors
201 required to feasibly schedule the same set of tasks. As shown in the Table
202 1, RMNF has a approximation ratio as 2.67, while RMFF and RMBF have
203 better performance.

204 4. Global Scheduling

205 In this section, we will dive into another branch of scheduling strategies
206 – global approaches to multiprocessor real-time scheduling.

207 4.1. Overview

208 Global scheduling algorithms, as its name suggests, globally schedules
209 any feasible periodic task set. In contrast to partitioned scheduling, global
210 scheduling schedules jobs and tasks in one single shared queue instead of
211 multiple local, dedicated queues. By this means, the automatic load balanc-
212 ing and lower average response time can be achieved by global approaches.
213 In addition to this, global approaches also take advantages of the simpler
214 implementations and the existence of optimal schedulers.

215 Global scheduling strategies include Global Fixed-Job-Priority Scheduling,
216 Global Fixed-Task-Priority Scheduling, and Global Dynamic Priority
217 Scheduling. Although there are various categories of global scheduling algo-
218 rithms, the focus of this paper is on the Global Dynamic Priority Scheduling.
219 In the following subsection, it will be characterized in details.

220 4.2. Global Dynamic Priority Scheduling

221 In this subsection, we will present our in-depth exploration to the track
222 of global dynamic priority scheduling algorithm. To the best of our knowl-
223 edge, a number of global dynamic priority scheduling algorithms are optimal
224 for periodic tasksets with explicit or implicit deadlines. For example, Pro-
225 portionate Fairness algorithm and its variants including PD, PD², ERFair,
226 BF, SA [9], and LLREF [10] as well, are all optimal for offline environment.
227 Nevertheless, no algorithms until now are optimal to cope with online pre-
228 emptive scheduling problem, where tasksets are sporadic and multi-processor
229 environments are enforced. On the other hand, despite of its optimality and
230 dominance in theory, the usage of global dynamic priority algorithms are
231 limited in practice. This is because the existence of frequent preemption and
232 migration between tasks gives rise to excessive overheads in potential.

233 The following part of this subsection will provide brief summary of three
234 classic global dynamic priority scheduling algorithms. They are respectively
235 Proportionate Fairness Algorithm (PFair), and Largest Local Remaining Ex-
236 ecution First (LLREF).

237 4.2.1. *PFair*

238 Baruah et al [11] introduced Proportionate Fairness Algorithm. The Pfair
 239 class of algorithms that allow full migration and fully dynamic priorities have
 240 been shown to be theoretically optimal – i.e., they achieve a schedulable
 241 utilization bound (below which all tasks meet their deadlines) that equals
 242 the total capacity of all processors. Here are some fundamental principles of
 243 Proportionate Fairness algorithms:

- 244 1. Timeline is divided into equal length slots.
- 245 2. Task period and execution time are multiples of the slot size.
- 246 3. Each task receives amount of slots proportional to its task utilization.

247 The essential part of PFair algorithms is the quantum-based optimization
 248 defined over the lag of each task $lag(\tau_i, t)$, with the goal of minimizing the
 249 maximum lags of all tasks $\max_t |lag(\tau_i, t)|$.

$$\underbrace{lag(\tau_i, t)}_{error} = \underbrace{t \cdot \left(\frac{c_i}{T_i}\right)}_{fluid\ execution\ in[0,t)} - \underbrace{allocated(\tau_i, t)}_{real\ execution\ in[0,t)} \quad (4)$$

250 The generation of an optimal schedule is based on the above definition of
 251 lag . PFair algorithm does execute all urgent tasks with $lag(\tau_i, t) > 0$ and
 252 $lag(\tau_i, t + 1) \geq 0$ if τ_i executes. On top of that, PFair algorithm does not
 253 execute tnegru tasks, for which $lag(\tau_i, t) < 0$ and $lag(\tau_i, t + 1) \leq 0$ if τ_i does
 254 not execute. Besides, for other tasks, only those that have the least t such
 255 that $lag(\tau_i, t) > 0$ are executed.

256 The PFair algorithm will assign priorities to tasks at every time slot,
 257 which indicated itself as one of job-level dynamic priority scheduling policies.
 258 However, this characteristic gives rise to some issues, for example, frequent
 259 preemptions and frequent migrations.

260 Proportionate Fairness algorithm, as a strong candidate for solving re-
 261 source allocation problems, has wide variety of interesting applications and
 262 powerful theoretical support. For instances, Kelly et al [12] presented its
 263 application on the problem of rate control for communication networks. Ap-
 264 plication of PFair algorithm on LANs and hoc networks was fulfilled by Jiang
 265 et al [13]. Besides, Bonald's paper [14] provides in-depth queueing analysis
 266 over proportionate fairness as well as max-min fairness and balanced fairness.

267 4.2.2. *LLREF*

268 LLREF was firstly introduced by Cho et al [10]. Similar to PFair class
269 of algorithms, LLREF is also based on the fluid scheduling model, where
270 each task executes at a constant rate at all times. The principal idea of
271 LLREF is that given M processors, M largest local remaining execution
272 time tasks are selected first for every secondary event. This is also called
273 the LLREF scheduling policy. To reason over task execution behavior on
274 multiprocessors, a novel abstraction called Time and Local Execution Time
275 Domain Plane (T-L Plane) was developed.

276 This algorithm divides the schedule into Time and Local execution time
277 planes (TL-planes), which are determined by task deadlines. The algorithm
278 schedules tasks by creating smaller local jobs within each TL-plane. The
279 only parameters considered by the algorithm during a TL-plane are the pa-
280 rameters of the local jobs. When a TL-plane completes, the next TL-plane
281 is started. The duration of each TL-plane is the amount of time between
282 consecutive deadlines.

283 4.3. *Summary*

284 The global scheduling paradigm has advantages over the partitioned ap-
285 proach. First of all, if tasks can join and leave the system at run-time,
286 then it may be necessary to reallocate tasks to processors in the partitioned
287 approach. In addition, the partitioned approach cannot produce optimal
288 real-time schedules – one that meets all task deadlines when task utiliza-
289 tion demand does not exceed the total processor capacity – for periodic task
290 sets, since the partitioning problem is analogous to the bin-packing problem
291 which is known to be NP-hard in the strong sense. On top of that, in some
292 embedded processor architectures with no cache and simpler structures, the
293 overhead of migration has a lower impact on the performance. Finally, global
294 scheduling can theoretically contribute to an increased understanding of the
295 properties and behaviors of real-time scheduling algorithms for multiproces-
296 sors.

297 However, the global scheduling paradigm has also several disadvantages.
298 Firstly, global scheduling strategies are much more complicated to implement
299 than partitioned scheduling. In other words, for the partitioned approach,
300 once a set of tasks are allocated to processors, the multiprocessor real-time
301 scheduling problem becomes a collection of single processor real-time schedul-
302 ing problems. The ease of programming partitioned scheduling is obvious
303 since the single processor scheduling problem has already been well-studied

304 and optimal algorithms with easy implementations already exist. Secondly,
305 migrating tasks at run-time means more runtime overhead in that migrating
306 tasks may suffer cache misses on the newly assigned processor. If the task set
307 is fixed and known in advanced, it is obvious that the partitioned approach
308 provides more appropriate solutions.

309 **5. Hybrid Approaches**

310 **6. Conclusions**

311 References

- 312 [1] C. L. Liu, J. W. Layland, Scheduling algorithms for multiprogramming
313 in a hard-real-time environment, *Journal of the ACM (JACM)* 20 (1973)
314 46–61.
- 315 [2] L. Sha, R. Rajkumar, J. P. Lehoczky, Priority inheritance protocols: An
316 approach to real-time synchronization, *Computers, IEEE Transactions*
317 on 39 (1990) 1175–1185.
- 318 [3] A. Mohammadi, S. G. Akl, Scheduling algorithms for real-time systems,
319 School of Computing, Queens University (2005).
- 320 [4] R. I. Davis, A. Burns, A survey of hard real-time scheduling for multi-
321 processor systems, *ACM Computing Surveys (CSUR)* 43 (2011) 35.
- 322 [5] M. R. Garey, D. S. Johnson, *Computers and intractability: a guide*
323 *to the theory of np-completeness*. 1979, San Francisco, LA: Freeman
324 (1979).
- 325 [6] J. Carpenter, S. Funk, P. Holman, A. Srinivasan, J. Anderson,
326 S. Baruah, A categorization of real-time multiprocessor scheduling prob-
327 lems and algorithms, *Handbook on scheduling algorithms, methods, and*
328 *models* (2004) 30–1.
- 329 [7] S. K. Dhall, C. Liu, On a real-time scheduling problem, *Operations*
330 *research* 26 (1978) 127–140.
- 331 [8] Y. Oh, S. H. Son, Tight performance bounds of heuristics for a real-time
332 scheduling problem (1993).
- 333 [9] A. Khemka, R. Shyamasundar, An optimal multiprocessor real-time
334 scheduling algorithm, *Journal of parallel and distributed computing* 43
335 (1997) 37–45.
- 336 [10] H. Cho, B. Ravindran, E. D. Jensen, An optimal real-time scheduling
337 algorithm for multiprocessors, in: *Real-Time Systems Symposium, 2006.*
338 *RTSS’06. 27th IEEE International, IEEE*, pp. 101–110.
- 339 [11] S. K. Baruah, N. K. Cohen, C. G. Plaxton, D. A. Varvel, Proportionate
340 progress: A notion of fairness in resource allocation, *Algorithmica* 15
341 (1996) 600–625.

- 342 [12] F. P. Kelly, A. K. Maulloo, D. K. Tan, Rate control for communication
343 networks: shadow prices, proportional fairness and stability, Journal of
344 the Operational Research society (1998) 237–252.
- 345 [13] L. B. Jiang, S. C. Liew, Proportional fairness in wireless lans and ad hoc
346 networks, in: Wireless Communications and Networking Conference,
347 2005 IEEE, volume 3, IEEE, pp. 1551–1556.
- 348 [14] T. Bonald, L. Massoulié, A. Proutiere, J. Virtamo, A queueing anal-
349 ysis of max-min fairness, proportional fairness and balanced fairness,
350 Queueing systems 53 (2006) 65–84.